

No Tux Given

Diving Into Contemporary
Linux Kernel Exploitation



About Me

- Sam (@sam4k1)
- Background in VR and exploit dev
- I like Linux, security, games & food

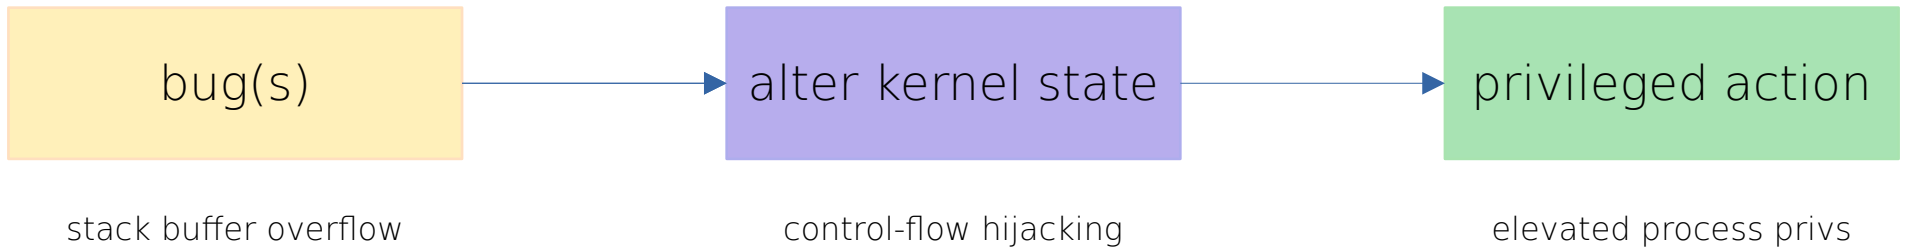
What Are We Doing Here?

- Exploring the past, present & future of kernel security & xdev
- Hopefully making an increasingly complex topic more accessible
- Do we need any more reasons??? This stuff is awesome!

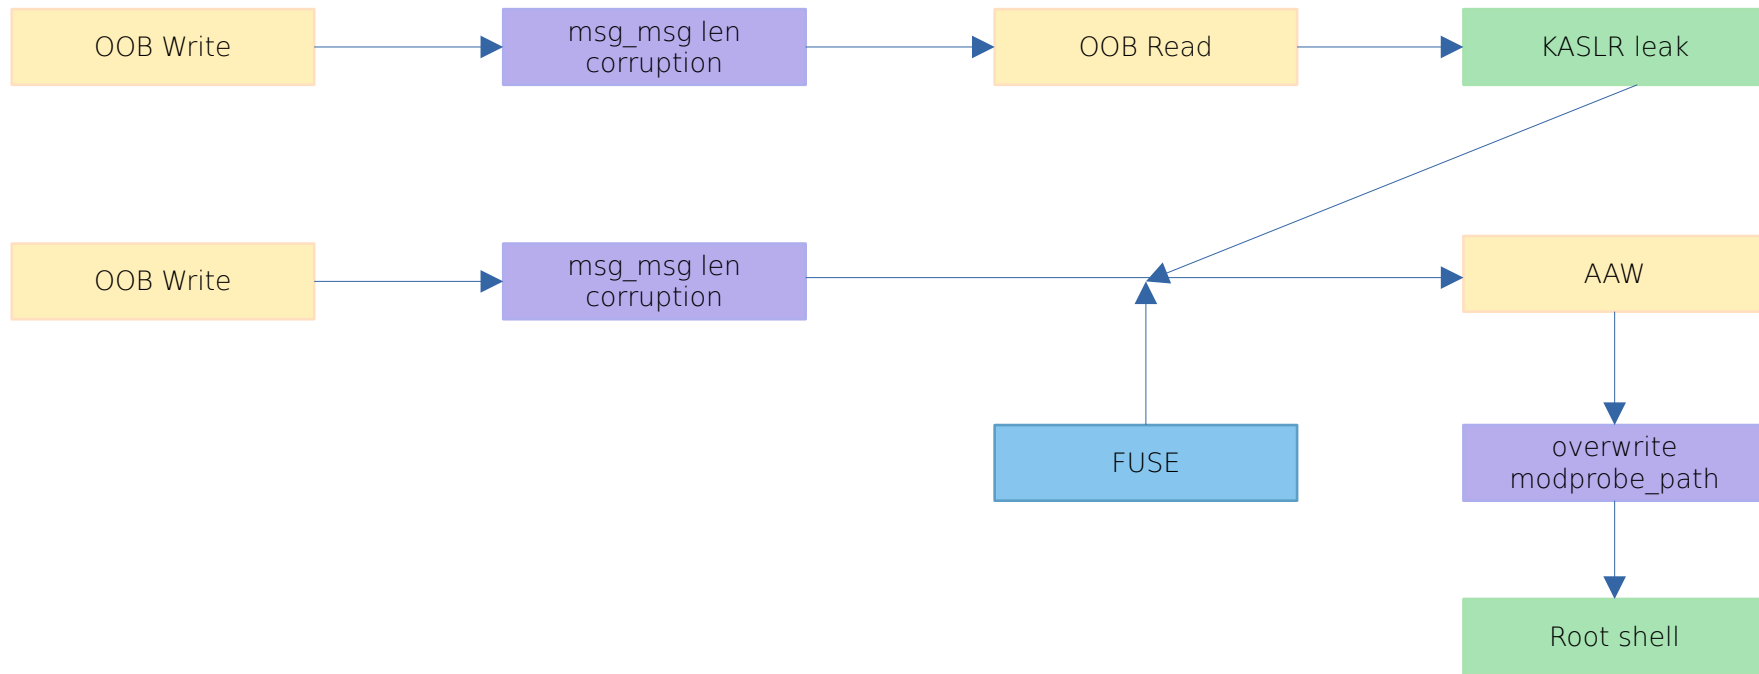
Tl;dr kernel exploits??



Tl;dr kernel exploits??




TL;dr kernel exploits??




Tux's Security Past

Examining Historical Kernel Exploitation Trends


Mitigations



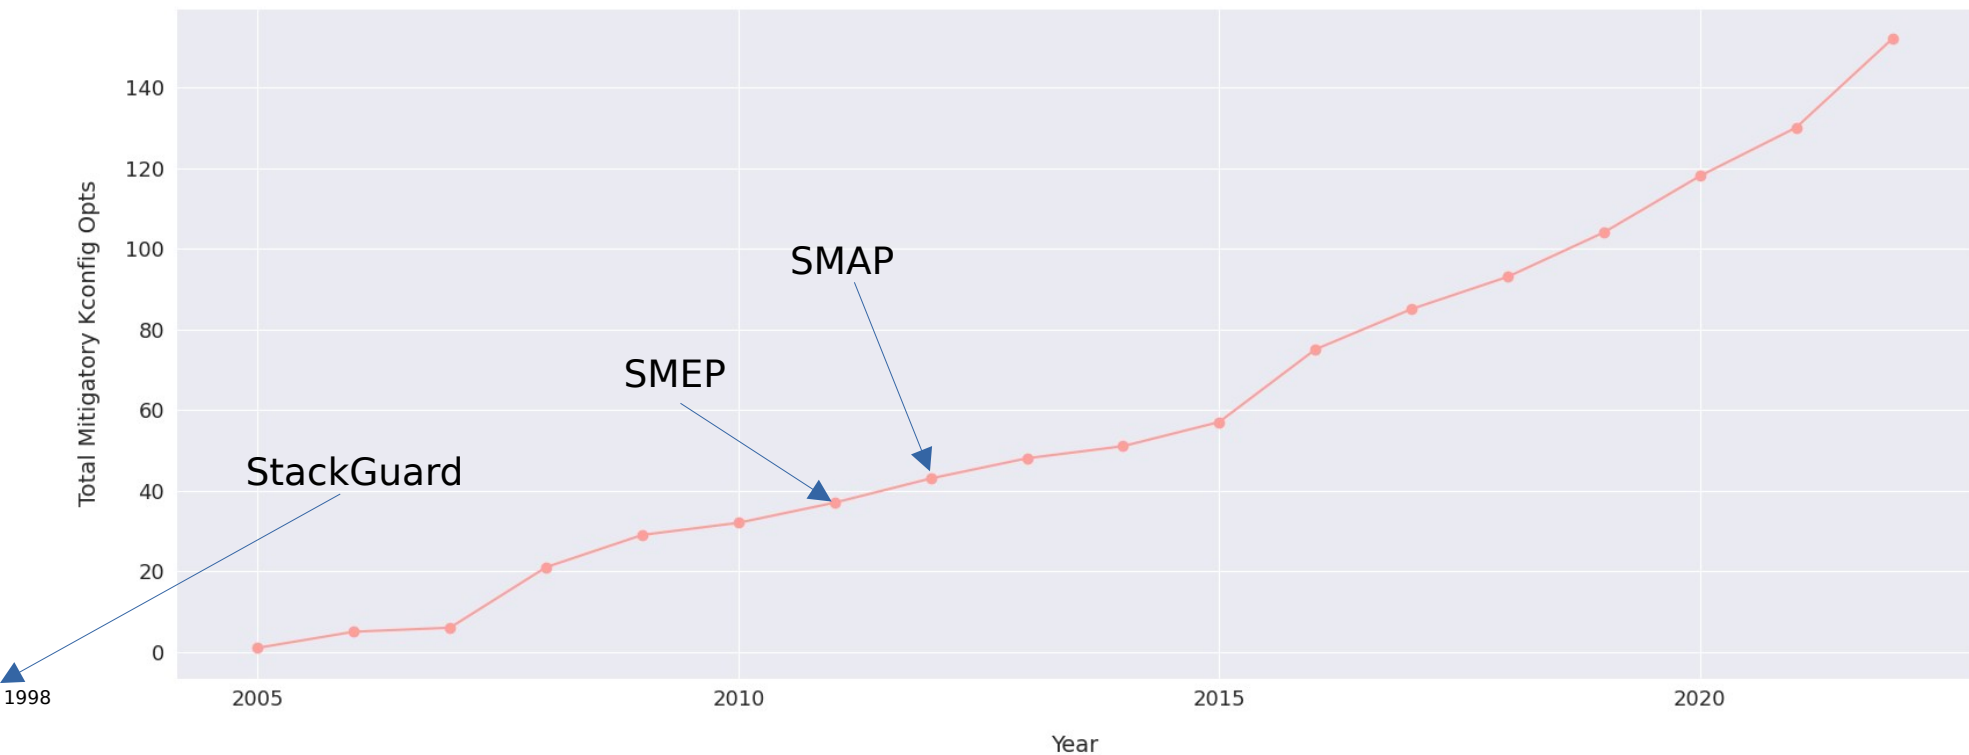
Mitigations




Mitigations




Mitigations




Mitigations




Mitigations



Bug Trends

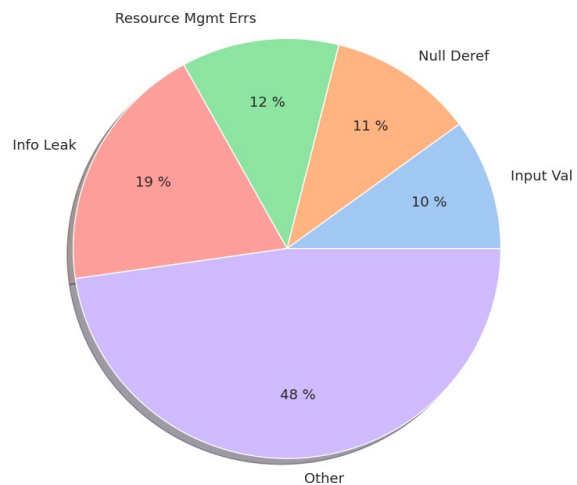


Bug Trends




Bug Trends


Top CWEs 2010



Top CWEs 2015



Top CWEs 2020



Tux's Security Present

Looking At Contemporary Kernel Exploitation

Kernel Exploits in 2023 | The Process

The process of getting from bug to privesc has become more complex:

- 1) Need to understand the attack surface
- 2) Find yourself some bugs (ezpz right?)
- 3) Figure out how, and what you need, to exploit it
 - Typically takes knowledge of platform/surface/bug and existing techniques
- 4) Actually get a (reliably??) working proof-of-concept

Kernel Exploits in 2023 | The Mindset

- Curiosity! Ask questions and take the time to understand
- Patience helps too, as sometimes there are no solutions
- Document, document, document! You'll thank yourself
- Opt for generic tooling and techniques where possible, to reuse
- The kernel is unforgiving of mistakes and unexpected behaviour!

Understanding The Attack Surface

- Informs where to look for bugs, what to look for and how to exploit them
- Lots of factors to consider: Kconfig, arch, platform specifics, 3rd parties etc.
- Varies greatly across desktop, android, IoT



Finding Some Bugs | Approaches

- Doesn't have to be 0days! Syzbot dashboard, silent fixes, n-days etc.
- QEMU + gdb make it easy to dig deeper and do some dynamic analysis
- Time spent understanding the bug & surface will help going forward
- Factor in surface/mitigations when thinking about what to look for

Finding Some Bugs | Tools & Tips

But if you do want a shiny 0day there's...

- Good ol' fashioned code auditing
- CodeQL to help flag areas of interest or check for specific patterns
- Spin up your own modified syzkaller instance
 - Adding coverage for areas without descriptions (e.g. 3rd party drivers)
 - Extending coverage for more tailored fuzzing using platform knowledge

From Bug To

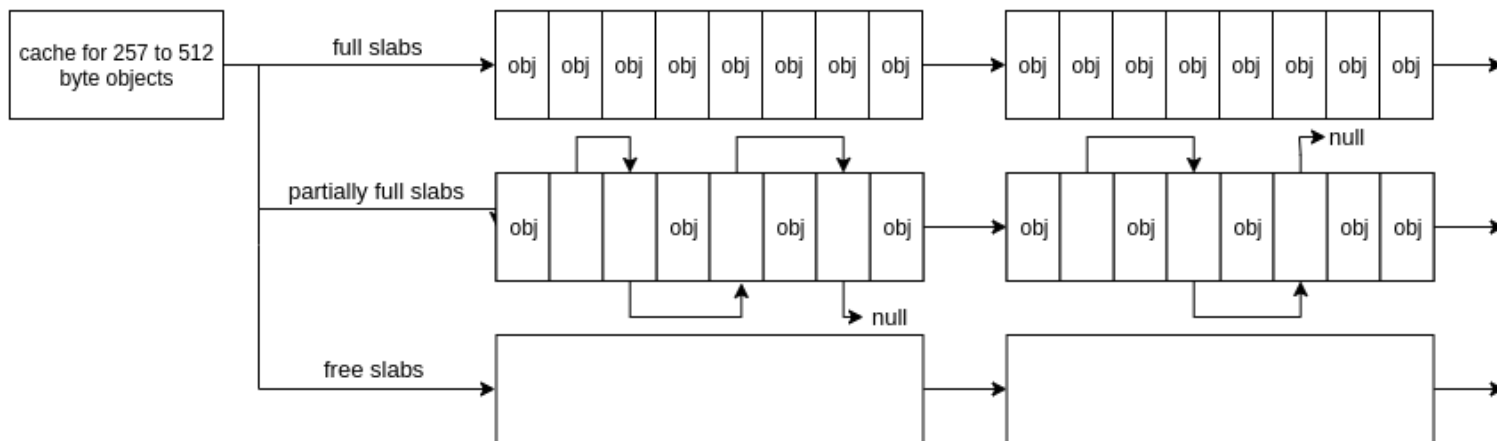
- Bug provides our initial primitive
- Generic techniques & strategies to leveraging particular primitives
- With each surface/bug often having its own nuances & requirements
- Goal is to chain these together to ultimately privesc
 - Typically via elevating our procs privs or executing another bin with privs

Exploiting UAFs | Getting Our Bearings

- Can cause the kernel to do some action(s) on previously freed memory
- So we need to think about how the kernel allocates this memory:

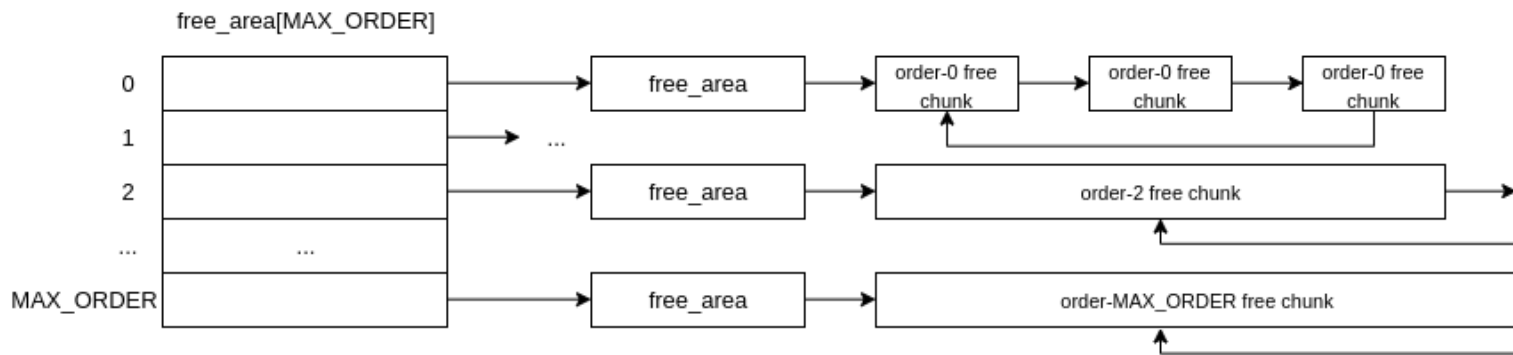
Exploiting UAFs | Getting Our Bearings

- Can cause the kernel to do some action(s) on previously freed memory
- So we need to think about how the kernel allocates this memory:
 - SLUB allocator: used for small, commonly used objects



Exploiting UAFs | Getting Our Bearings

- Can cause the kernel to do some action(s) on previously freed memory
- So we need to think about how the kernel allocates this memory:
 - Page allocator: handles larger, contiguous allocs (including slabs!)



(Where chunk size = $2^{\text{order}} * \text{PAGE_SIZE}$)

Exploiting UAFs | Getting Our Bearings

- Can cause the kernel to do some action(s) on previously freed memory
- So we need to think about how the kernel allocates this memory:
 - SLUB allocator: used for small, commonly used objects
 - Page allocator: handles larger, contiguous allocs (including slabs!)
- We also need to consider what actions are done on the freed memory
- As well as how reachable/triggerable the UAF is and any timing issues

Exploiting UAFs | Mitigations

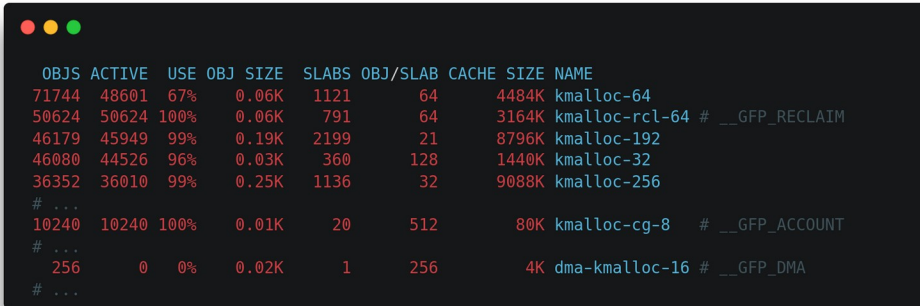
	Ubuntu 22.04 (5.15)	kCTF (6.1)	Pixel 7 (5.10)
<code>init_on_alloc</code>	default	default	default
<code>SLAB_FREELIST_RANDOM</code>	default	default	default
<code>SHUFFLE_PAGE_ALLOCATOR</code>	default	not set	default
<code>STATIC_USERMODEHELPER</code>	not set	not set	default
no unpriv userfaultfd OR FUSE	FUSE	neither enabled	neither unpriv*
<code>slab_nomerge</code>	not set	default	default

Exploiting UAFs | Realising Our Goal

- Need an object to replace our freed one
 - Such that actions done on it give us further kernel primitives/priv esc
 - CodeQL is a useful tool here to query specific obj criteria (size, offsets etc.)
- Then we need to make sure our object(s) ends up where its supposed to...
 - i.e. we need to understand how to control the memory layout

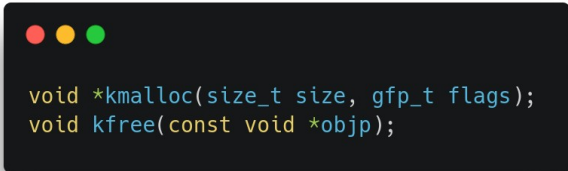
Exploiting UAFs | Shaping General Purpose Caches

- Different **gfp_t flags** may be allocated into different general purpose caches
 - E.g. **GFP_KERNEL_ACCOUNT**, used for objects containing user data
- Elastic objects provide us with a generic approach, usable across cache sizes
- Cache noise is also an important factor in tuning the reliability of your heap spray
- FUSE can open up more allocation possibilities by allowing us to keep more ephemeral object allocations in memory^[8]



OBJS	ACTIVE	USE	OBJ	SIZE	SLABS	OBJ/SLAB	CACHE	SIZE	NAME
71744	48601	67%	0.06K	1121	64	4484K	kmalloc-64		
50624	50624	100%	0.06K	791	64	3164K	kmalloc-rcl-64	# __GFP_RECLAIM	
46179	45949	99%	0.19K	2199	21	8796K	kmalloc-192		
46080	44526	96%	0.03K	360	128	1440K	kmalloc-32		
36352	36010	99%	0.25K	1136	32	9088K	kmalloc-256		
# ...									
10240	10240	100%	0.01K	20	512	80K	kmalloc-cg-8	# __GFP_ACCOUNT	
# ...									
256	0	0%	0.02K	1	256	4K	dma-kmalloc-16	# __GFP_DMA	
# ...									

\$ sudo slabtop



```
void *kmalloc(size_t size, gfp_t flags);  
void kfree(const void *objp);
```

API example for general purpose allocs

Exploiting UAFs | Shaping Private Caches

```
struct kmem_cache *kmem_cache_create(const char *name, unsigned int size,
                                     unsigned int align, slab_flags_t flags,
                                     void (*ctor)(void *));
void *kmem_cache_alloc(struct kmem_cache *s, gfp_t flags);
void kmem_cache_free(struct kmem_cache *s, void *objp);
```

API example for private cache allocs

```
  OBUS ACTIVE  USE OBJ SIZE  SLABS OBJ/SLAB  CACHE SIZE  NAME
754992 754536 99%  0.19K 35952    21    143808K dentry
493056 490561 99%  0.03K  3852    128    15408K numa_policy
343280 342993 99%  0.57K 12260    28    196160K radix_tree_node
272944 272913 99%  1.12K  9748    28    311936K btrfs_inode
269640 267748 99%  0.07K  4815    56    19260K vmap_area
265216 134497 50%  0.01K   518   512    2072K lsm_file_cache
206808 198880 96%  0.14K  7386    28    29544K btrfs_extent_map
138560 117011 84%  0.06K  2165    64    8660K dmaengine-unmap-2
123100 122135 99%  0.16K  4924    25    19696K vm_area_struct
 74752  73535 98%  0.06K  1168    64    4672K anon_vma_chain
```

\$ sudo slabtop

- Same goal as before, except...
- These caches only contain specified obj
- But... the slabs that make up private and general purpose caches are allocated the same way, by the buddy allocator
- With a bit more work, tuning and luck it's possible to have the slab containing freed private obj to be reallocated as a slab for a general purpose cache
- AKA cross-cache attacks

Exploiting UAFs | Shaping The Buddy (Page) Allocator

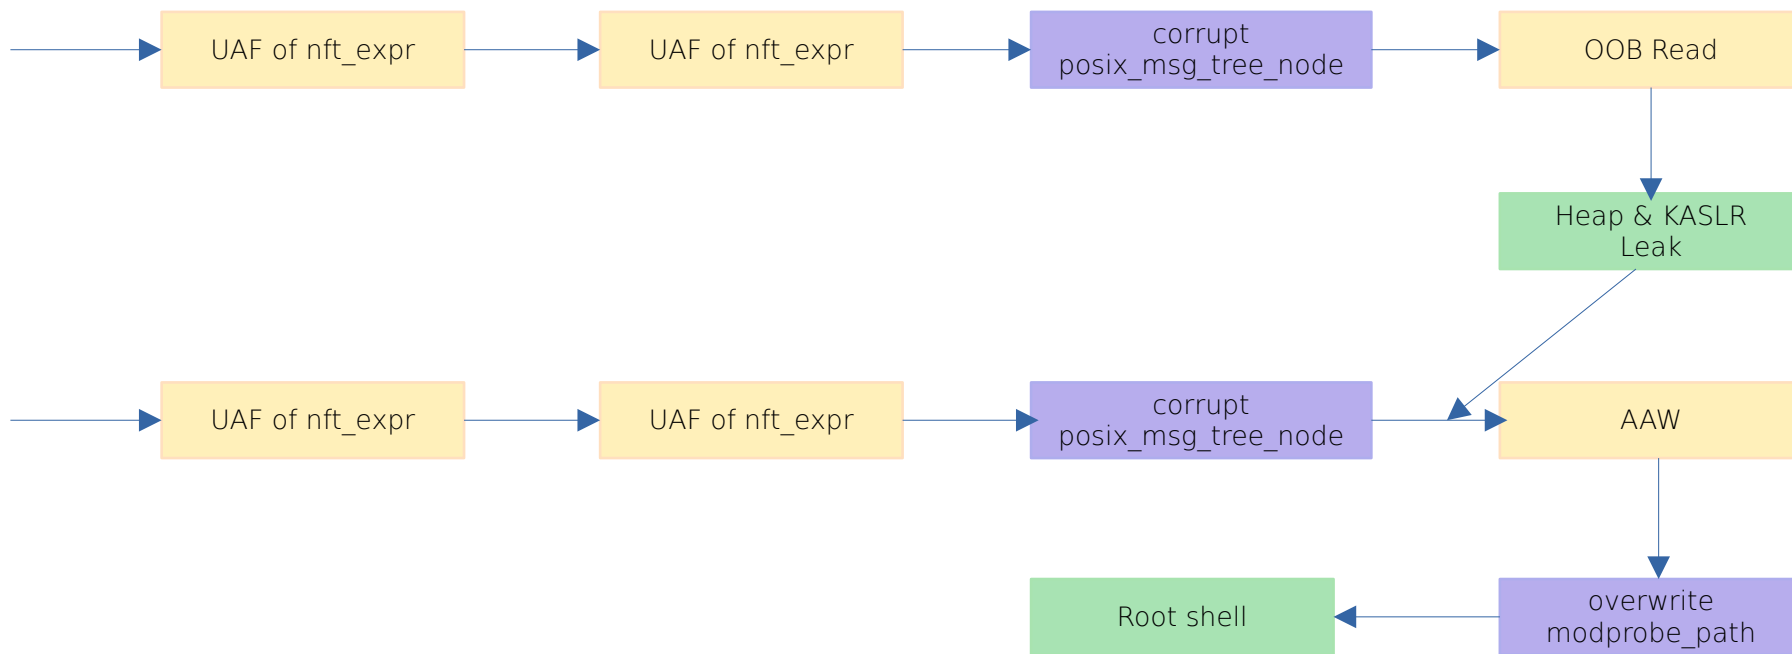
- Goal is the same, just need to remember the different structure!
- Used for large dynamic buffers (GPU, packet ring buffers), *slabs*
- Need to mitigate noise from chunks merging
 - If lower order is empty, chunks are split
 - If higher order is empty, *contiguous* chunks merged
- May also want to ensure contiguity of multiple allocations

```
struct page *alloc_pages(gfp_t gfp, unsigned int order);
unsigned long __get_free_pages(gfp_t gfp_mask, unsigned int order);
void __free_pages(struct page *page, unsigned int order);
void free_pages(unsigned long addr, unsigned int order);
static void *__kmalloc_large_node(size_t size, gfp_t flags, int node);
```

```
----- zone info -----| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10
Node 0, zone DMA 1 1 1 1 1 1 1 1 1 0 1 2
Node 0, zone DMA32 37337 28593 15203 5885 1783 703 322 208 145 147 198
Node 0, zone Normal 4553 1541 3541 1734 952 396 159 78 29 16 107
```


Exploiting UAFs | A Real World Example

- CVE-2022-32250^[5] was a UAF in Netfilter:



Exploiting OOB Writes | What Bounds?

- Different kinds out-of-bounds writes in the kernel...
 - Array indexes, heap overflows, stack overflows etc.
- However this list may be shorter after we factor in mitigations...

Exploiting OOB Writes | Mitigations

	Ubuntu 22.04 (5.15)	kCTF (6.1)	Pixel 7 (5.10)
FORTIFY_SOURCE	default	default	default
UBSAN	UBSAN_TRAP not set	not set	default
SLAB_FREELIST_HARDENED	default	not set	default
STATIC_USERMODEHELPER	not set	not set	default

Exploiting OOB Writes | Heap Overflows It Is

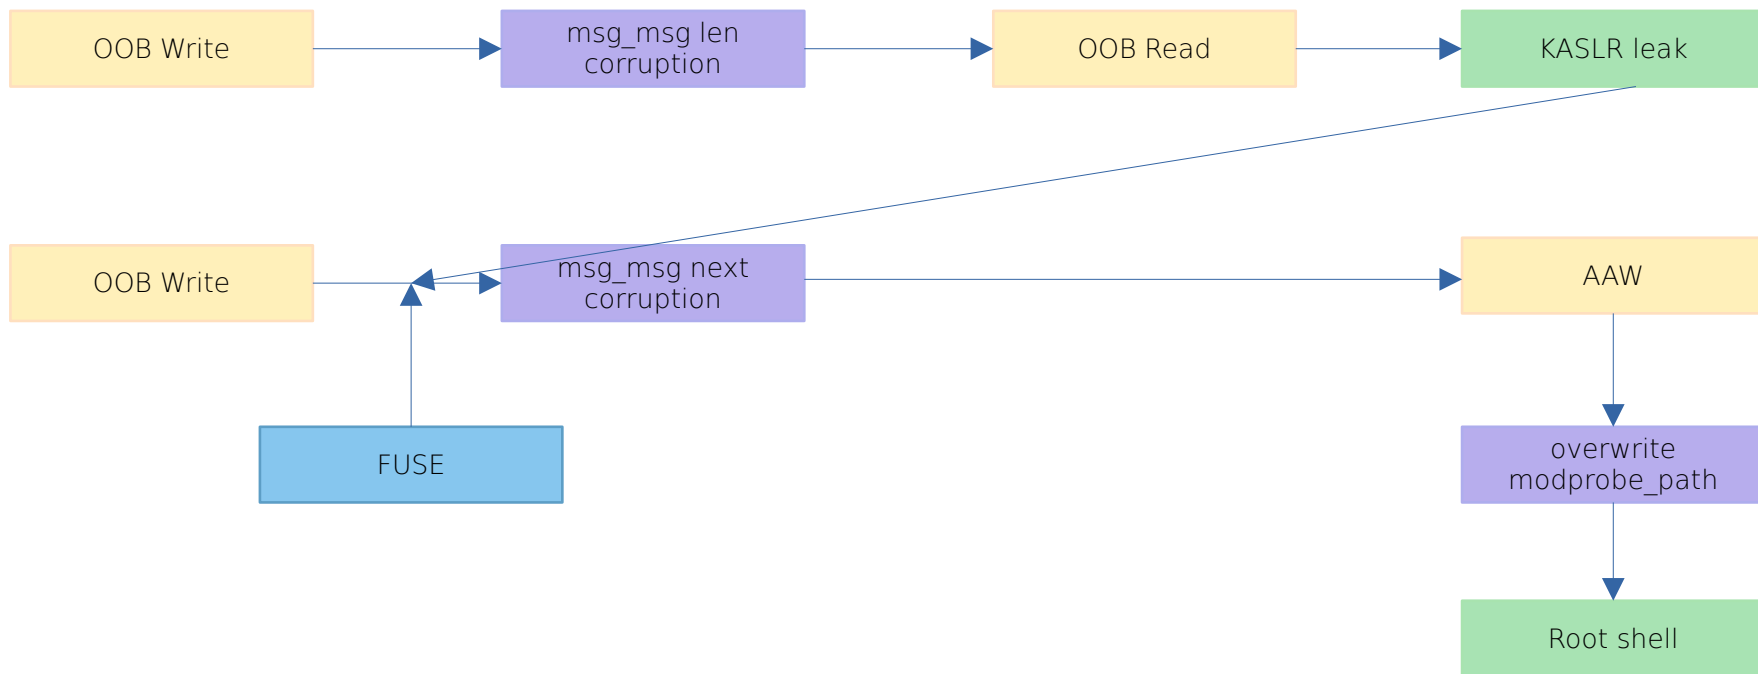
- As we're dealing with the heap: how is our object allocated?
- Now interested in what's adjacent to our object:
 - Another object we can corrupt?
 - The freelist pointer for the slab?
 - Another buddy allocated chunk?
- What is the extent of our overflow? Controlled size/data?
- With all this info, we can find a suitable candidate to corrupt

Exploiting OOB Writes | Getting To The Finish Line

- Want to pivot from our initial OOB write primitive
- Elastic objects are a popular target
 - E.g. **msg_msg** can be used to pivot from an OOB write to AAW^[6]
- Cross-cache attacks open up possible targets to sensitive, otherwise inaccessible corruption targets
- **modprobe_path** is still an easy target to privesc with an AAW

Exploiting OOB Writes | A Real World Example

- CVE-2022-0185^[7] was a heap overflow in fsconfig(2):



Exploiting Race Conditions

- Typically enable other bugs, such as use-after-frees
- But can be hard to debug; how do we know we're even winning the race?!
 - Printk debugging ☺ (or other kernel instrumentation, e.g. sleeps to widen the race)
 - Gdb scripts can also make life easier here
- And if we can win it, what if the odds are super low?
 - FUSE (or userfaultfd on older systems) may be an option for hanging kernel execution
 - Alternatively, user-triggerable interrupts (e.g. timers) can widen race condition too^[11]
- Be considerate of the little gotchas
 - Execution contexts? Locks? Who's executing what, when? CPU affinity? etc.

Tux's Security Future

Some Thoughts On Future Impacts to Kernel Security

Looking Ahead



Looking Ahead | New Mitigations

- kCTF experimental mitigations^[10]:
 - **KMALLOC_SPLIT_VARSIZE**: mitigate generic direct object reuse via elastic objects (looking at you **msg_msg**!)
 - **SLAB_VIRTUAL**: mitigate cross-cache attacks by reworking slab mem use
- Worth noting that many proprietary mitigations don't yet have mainline equivalents (e.g. grsec's **AUTOSLAB**)
- Lag between mainline mitigation support & hardware adoption
 - E.g. Intel's CFI (CET) support was introduced in their 11th Gen CPUs (2021)

Looking Ahead | New Technologies (AKA Rust)

- Yep, it's Rust time
- Initial support released in kernel version 6.1
- Memory safety built-in as opposed to being bolted on
- Where 66% of kernel security issues are memory safety related (2019)^[9]
- However, Rust is still a tool used by people, and we make mistakes!

Looking Ahead | Attitude to Security

- Finding the balance between performance/usability and security
 - When to include, and default, particular mitigations?
 - Most of the topics mentioned today have mitigations
- Fostering open and accessible environment for security research
 - Public research and sharing can drive innovation and improvements
 - Vs. malicious actors who are happy to keep all this in the shadows
 - Still friction in the handling of security fixes & disclosures



Wrapping Up

Thank You! Feel Free To @ Me Online/Offline

Resources

- <https://github.com/xairy/linux-kernel-exploitation>
- <https://github.com/a13xp0p0v/linux-kernel-defence-map>
- <https://sam4k.com> (any talk updates will be posted here!)
- <https://codeql.github.com>
- <https://github.com/google/syzkaller>

Refs

1. <https://github.com/a13xp0p0v/kconfig-hardened-check/>
2. <https://cateee.net/lkddb/web-lkddb/>
3. <https://github.com/cloudsecurityalliance/gsd-database>
4. <https://github.com/torvalds/linux>
5. <https://blog.theori.io/research/CVE-2022-32250-linux-kernel-lpe-2022/>
6. <https://www.willsroot.io/2021/08/corctf-2021-fire-of-salvation-writeup.html>
7. <https://www.willsroot.io/2022/01/cve-2022-0185.html>
8. <https://exploiter.dev/blog/2022/FUSE-exploit.html>
9. https://static.sched.com/hosted_files/lssna19/d6/kernel-modules-in-rust-lssna2019.pdf
10. https://github.com/thejh/linux/blob/slub-virtual-v6.1-lts/MITIGATION_README
11. <https://googleprojectzero.blogspot.com/2022/03/racing-against-clock-hitting-tiny.html>